



POLITECNICO
MILANO 1863

GPUs and Heterogeneous Systems
(programming models and architectures)

Performance considerations

Optimizing the performance of CUDA kernels

- Accelerating a function in CUDA is pretty simple
- Optimizing the performance of the accelerated function is a hard task
 - It is necessary to “tailor” and tune the kernel code based on the underlying architecture
 - This is a profile-driven activity
 - Identify the bottleneck and work on it!
- This presentation presents an overview of possible optimizations

Maximizing occupancy

- **Benefit to compute cores**
 - More work to hide long-latency operations
- **Benefit to memory**
 - More parallel memory accesses to hide memory latency
- **Strategies**
 - Tuning usage of streaming multiprocessors' resources (threads per block, shared memory per block, registers per thread)

Enabling coalesced global memory accesses

- **Benefit to compute cores**
 - Fewer stalls waiting for global memory accesses
- **Benefit to memory**
 - Less global memory traffic and better utilization of caches
- **Strategies**
 - Transferring data from global memory to shared memory in a coalesced way and accessing shared memory in a uncoalesced way
 - Rearranging thread to data mapping
 - Rearranging the layout of the data

Minimizing control divergence

- **Benefit to compute cores**
 - High SIMT efficiency
- **Benefit to memory**
 - N/A
- **Strategies**
 - Rearranging thread to data/work mapping
 - Rearranging the layout of the data

Tiling of reused data

- **Benefit to compute cores**
 - Fewer stalls waiting for global memory accesses
- **Benefit to memory**
 - Less global memory traffic
- **Strategies**
 - Placing data that is reused within a block in shared memory or registers so that it is transferred from global memory only once

Privatization

- **Benefit to compute cores**
 - Fewer stalls waiting for atomic updates
- **Benefit to memory**
 - Less contention and serialization of atomic updates
- **Strategies**
 - Applying partial updates to a private copy of the data and then updating the final result

Thread coarsening

- **Benefit to compute cores**
 - Less redundant work, divergence, or synchronization
- **Benefit to memory**
 - Less global memory traffic
- **Strategies**
 - Assigning multiple chunks of data to be processed to each thread to reduce the price of parallelism when it is incurred unnecessary

Other instruction-level / architecture specific optimizations

- Properly selecting data types
- Use intrinsic functions offering higher performance at a reduced result accuracy
- Follow guidelines for the specific GPU family

References

- Slides mainly based on:
 - W.-m. W. Hwu , D. B. Kirk, I. El Hajj, **Programming Massively Parallel Processors: A Hands-on Approach, Chapter 6**
 - J. Chen, M. Grossman, T. McKercher, **Professional Cuda C Programming, Chapters 3, 4, 5, 7**