

GPUs and Heterogeneous Systems – A.Y. 2023-24

Scuola di Ingegneria Industriale e dell'Informazione
Prof. Antonio Miele



POLITECNICO
MILANO 1863

July 19, 2024 - **FIRST PART OF THE EXAM**

Surname:	Name:	Personal Code:
----------	-------	----------------

Question	1	2	3	4	5	OVERALL
Max score	3	3	3	3	3	15
Score						

Instructions:

- This first part of the exam is “closed book”. The students are not allowed to consult any course material and notes.
- No extra devices (e.g., phones, iPad) are allowed. Please, shut down and store any electronic device.
- Students are not allowed to communicate with any other ones.
- Students can write in pen or pencil, any color, but avoid writing in red.
- Any violation of the above rules will lead to the invalidation of the test.
- **Duration: 30 minutes**

Question 1

Briefly explain what the graphics pipeline is.

Question 2

Complete the following OpenACC pragmas to optimize data transfer between the host and the device. In particular, 1) specify between brackets () the appropriate list of arrays (index range for each array is always [0:N]), and 2) delete unnecessary clauses. Motivate your answer.

```
void foo(int *A, int *B){
    int C[N];
    #pragma acc data copy(    ) copyin(    ) copyout(    ) create(    )
    {
        #pragma acc parallel loop copy(    ) copyin(    ) copyout(    ) create(    )
        for(int i = 0; i < N; i++)
            C[i] = A[i] + B[i];
        #pragma acc parallel loop copy(    ) copyin(    ) copyout(    ) create(    )
        for(int i = 0; i < N; i++)
            A[i] = C[i] * A[i];
    }
}
```

Question 3

Simulate the following parallel reduction kernel (based on the algorithm optimizing memory access efficiency) and show the data array's content at each loop iteration. Consider a grid with 2 blocks, each block with 4 threads; the initial content of the data array is: [0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3]

```
#define STRIDE_FACTOR 2

__global__ void reduce(double* data) {
    int i = threadIdx.x;

    int base_i = blockDim.x * blockIdx.x * STRIDE_FACTOR;
    for (int stride = blockDim.x; stride >= 1; stride /= STRIDE_FACTOR) {
        if (i < stride) {
            data[base_i + i] += data[base_i + i + stride];
        }
        __syncthreads();
    }
}
```

Question 4

Draw a basic example of the roofline model. Then, add to the plot the following 3 points representing:

- A. A memory-bound application efficiently using resources
- B. A memory-bound application not efficiently using resources
- C. A compute-bound application efficiently using resources

Question 5

Briefly describe the CUDA memory model (organize the answer in a table); for each component, specify the name, which data is stored, type of access (read/write or read-only), and scope.

GPUs and Heterogeneous Systems – A.Y. 2023-24

Scuola di Ingegneria Industriale e dell'Informazione
Prof. Antonio Miele



POLITECNICO
MILANO 1863

July 19, 2024 - **SECOND PART OF THE EXAM**

Surname:	Name:	Personal Code:
----------	-------	----------------

Question	1	2	3	OVERALL
Max score	5.5	5.5	5	16
Score				

Instructions:

- This second part of the exam is “open book”. The students are allowed to use any material and notes.
- The students are allowed to use the laptop and the tablet. No extra devices (e.g., phones) are allowed. Please, shut down and store not allowed electronic devices.
- Students are not allowed to communicate with any other one or use Internet.
- Students can write in pen or pencil, any color, but avoid writing in red.
- Students can also use the laptop to code the test solution. In this case, please pay attention to the instructor’s instructions to submit the test solution.
- Any violation of the above rules will lead to the invalidation of the test.
- **Duration: 1 hour and 15 minutes**

Question 1

Implement two basic CUDA kernel functions, called `initialize` and `simulateStep`, to accelerate the compute-intensive portions of code in the following C program. Consider a 2D block shape.

Question 2

Modify the `main` function to execute the two CUDA kernel functions defined in the former question. Set block size to 32x32.

Question 3

Implement a new CUDA kernel function to accelerate the first compute-intensive function (`initialize`) using thread coarsening on the `x` axis; set a coarsening factor equal to 8. Specify if any change has to be applied in the `main` function.

The source code can be downloaded from the course page on WeBeep

```

/*
* The program simulates the heat transfer in a 2D surface. The 2D surface is modeled by
* means of a 2D array called surf (linearized for the sake of simplicity); each position in
* the array surf[i][j] contains the temperature of a single portion of the surface.
* Heat transfer is simulated by a partial differential equation (not discussed here in
* details...). The equation is solved by iterating over nsteps subsequent time steps; at
* each time step n, the new temperature of each position is computed as a function of the
* current temperatures of the same position surf[i][j] and of the 4 neighbor positions
* surf[i-1][j], surf[i][j-1], surf[i+1][j] and surf[i][j+1]; results are saved in a new
* array surf1. surf1 and surf are then swapped to simulate subsequent time step (surf1
* contains the input of the subsequent step while surf is used to store the results).
* When computing the temperature of a position on the border of the surface, the
* non-existing neighbor positions are replaced with the temperature of the central position
* itself (i.e., when computing the temperature surf1[0][2], the neighbor position
* surf[-1][2] is replaced with surf[0][2]).
*/

#include <stdio.h>
#include <stdlib.h>

#define A 0.5f // Diffusion constant
#define DX 0.01f // Horizontal grid spacing
#define DY 0.01f // Vertical grid spacing
#define DX2 (DX*DX)
#define DY2 (DY*DY)
#define DT (DX2 * DY2 / (2.0f * A * (DX2 + DY2))) // Largest stable time step
#define SMALLSIZEPROBLEM 100

int getIndex(int i, int j, int width);

// Take in input 2D coordinates of a point in a matrix
// and translate in a 1D offset for the linearized matrix
int getIndex(int i, int j, int width) {
    return (i*width + j);
}

int main(int argc, char **argv) {
    float *surf, *surf1, *tmp;
    int i, j, nx, ny, nsteps, n;

    // Read arguments
    if(argc != 4 || atoi(argv[3])%2!=0){
        printf("Please specify sizes of the 2D surface and the number of time steps\n");
        printf("The number of time steps must be even\n");
        return 0;
    }
    nx = atoi(argv[1]);
    ny = atoi(argv[2]);
    nsteps = atoi(argv[3]);

    // Allocate memory for the two arrays
    surf = (float*) malloc(sizeof(float) * nx * ny);
    if(!surf){
        printf("Error: malloc failed\n");
        return 1;
    }
    surf1 = (float*) malloc(sizeof(float) * nx * ny);
    if(!surf1){
        printf("Error: malloc failed\n");
        return 1;
    }

    // Initialize the data with a pattern of disk of radius of 1/6 of the width
    float radius2 = (nx/6.0f) * (nx/6.0f);
    for (i = 0; i < nx; i++) {
        for (j = 0; j < ny; j++) {
            int index = getIndex(i, j, ny);
            // Distance of point i, j from the origin
            float ds2 = (i - nx/2) * (i - nx/2) + (j - ny/2)*(j - ny/2);
            if (ds2 < radius2)
                surf[index] = 65.0f;
            else
                surf[index] = 5.0f;
        }
    }
}

```

```

// Compute-intensive kernel for the heat simulation
int index;
float sij, simlj, sijml, siplj, sijpl;

// Simulate N time steps
for (n = 0; n < nsteps; n++) {
    // Go through the entire 2D surface
    for (i = 0; i < nx; i++) {
        for (int j = 0; j < ny; j++) {
            // Compute the heat transfer taking old temperature from surf
            // and saving new temperature in surf1
            index = getIndex(i, j, ny);
            sij = surf[index];
            if(i>0)
                simlj = surf[getIndex(i-1, j, ny)];
            else
                simlj = sij;
            if(j>0)
                sijml = surf[getIndex(i, j-1, ny)];
            else
                sijml = sij;
            if(i<nx-1)
                siplj = surf[getIndex(i+1, j, ny)];
            else
                siplj = sij;
            if(j<ny-1)
                sijpl = surf[getIndex(i, j+1, ny)];
            else
                sijpl = sij;
            surf1[index] = sij + A * DT *
                ( (simlj - 2.0*sij + siplj)/DX2 + (sijml - 2.0*sij + sijpl)/DY2 );
        }
    }
    // Swap the surf and surf1 pointers for the next time step
    // (it avoids copying all data from surf1 to surf)
    tmp = surf;
    surf = surf1;
    surf1 = tmp;
}

// Print final temperatures ...

// Free memory
free(surf);
free(surf1);

return 0;
}

```