


Scan

What is a Scan?

- Takes as input $[x_0, x_1, \dots, x_{n-1}]$
 - Based on the return value, it could be
 - **Inclusive** each element includes the effect of the corresponding input elements
 - **Exclusive** each element does not include the effect of the corresponding input elements
- $[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})]$
 where i is the identity $[i, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})]$
- Also called **Prefix Sum**
 - The computational complexity of the sequential algorithm is $O(N)$

Usage Examples

- Example of where the work performed by some parallel algorithms can have higher complexity
- Used a primitive algorithm for different sorting algorithms
 - Radix sort
 - Quicksort
- Used to perform regex
 - For example the `grep` command



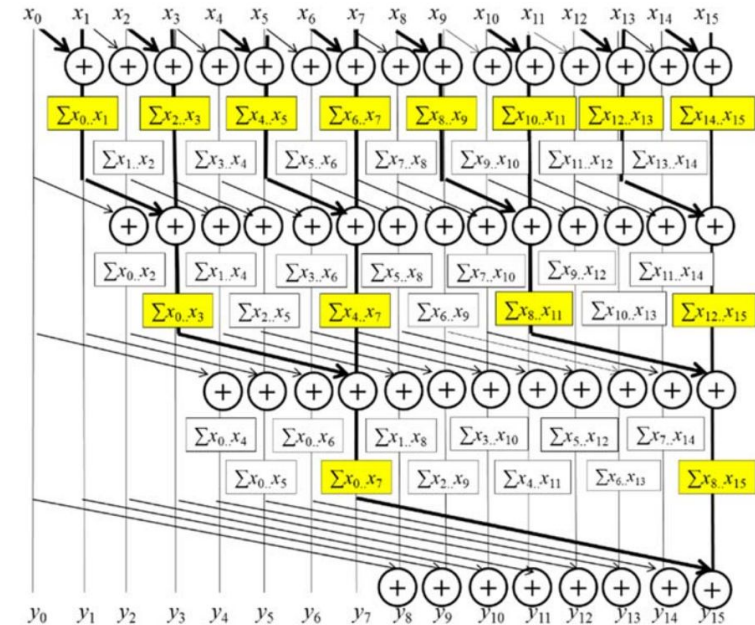
Naive Implementation

- Assumes that there are as many processors as data element
 - Each threads do the reduction for each output element
- Threads are doing redundant work
 - In the end, the bottleneck will be the longest path
 - The time required to get the last element
- It is not work efficient, the complexity is $O(N^2)$

$$\sum_{i=0}^{n-1} i = \frac{n \cdot (n - 1)}{2}$$


Kogge-Stone

- Before the algorithm begins, the `output[i]` contains the input element x_i .
- After k iterations, `output[i]` will contain the sum of the 2^k input elements before
- Write-after read hazards
 - Two `__syncthreads` are required
- The second `__syncthreads`
 - ensures that all active threads have completed their read of the old values
 - before any of them can move forward and perform a write.



Kogge-Stone: Double Buffering

- Write-after read hazards
- To avoid the second __syncthreads is to use double buffering
 - Each iteration you can swap the two
- It has a complexity of $O(N \log N)$



Work Efficiency

- It measures the extent to which the work performed is close to the minimum amount of work needed for the computation
- For example
 - The Kogge-Stone algorithm is not work efficient as the sequential one
 - The same applies to the naive one, which is also worse than the Kogge-Stone
- Kogge-Stone algorithm performs more computations
 - But it does so in fewer steps because of parallel execution
- With unlimited execution resources, the reduction would be approximately $N/\log(N)$

For $N=512$ would be $512/9=56.93x$



Code Hands-on

Bonus

- The [Brent-Kung](#) algorithm is another approach to this problem
- Performance-guided optimization in depth can be found [here](#)
- It is available in the [Thrust Library](#)
 - Documentation [here](#)



Thank you for your attention!

Gianmarco Accordi
gianmarco.accordi@polimi.it