

# Sparse Matrix Computation



# Sparse Matrices

---

- The majority of elements are zeros
  - Storing and processing these zero elements is a waste in terms of memory and bandwidth
- Different methods used to compact sparse matrices
  - COO
  - ELL
  - CSR
- They introduce irregularity in data representation
- Unfortunately, such irregularity can lead to
  - Underutilization of memory bandwidth
  - Control flow divergence
  - Load imbalance


# Why does it matter?

---

- Sparse matrices arise in many scientific, engineering, and financial modeling problems
  - Matrices can be used to represent the coefficients in a linear system of equations
  - Each row of the matrix represents one equation of the linear system
- Usually, scientific problem modeling involves many variables
  - Only a small number of them are relevant with non-zero coefficients
- Matrices are often used in solving linear systems of  $N$  equations of  $N$  variables in the form of  $AX + Y = 0$
- Graph computation
  - For example, large networks sparsely connected

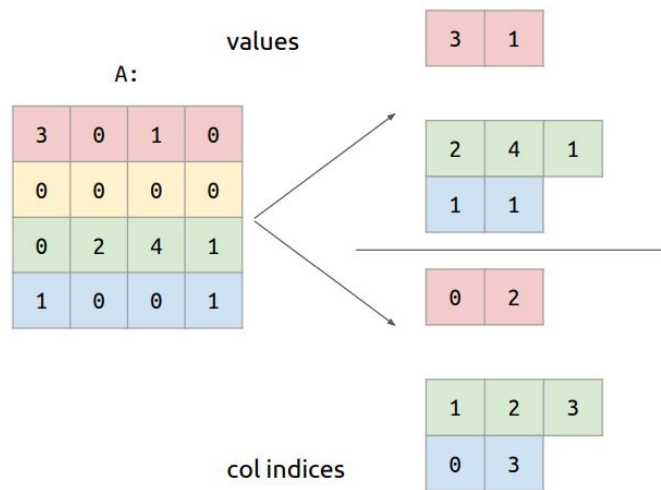
# SpMV

- **Sparse Matrix-Vector multiplication and accumulation (SpMV)**
  - Owing to its importance, standardized library function interfaces have been created to perform this operation
- SpMV is an **example** of the **tradeoffs** **between storage** formats
  - Space Efficiency
  - Flexibility
  - Accessibility
  - Memory Bandwidth
  - Load Balance
- Different sparse matrix storage formats remove all the zero elements from the matrix representation
  - No need to fetch them
  - No useless multiplications



# Compressed Sparse Row (CSR) Format

- Store each row as a sparse (row) vector
  - Each row is of variable length depending on the sparsity pattern



# Compressed Sparse Row (CSR) Format

- Additional storage is required to locate the start of each row

A:

3	0	1	0
0	0	0	0
0	2	4	1
1	0	0	1

values

3	1	2	4	1	1	1
---	---	---	---	---	---	---

col indices

0	2	1	2	3	0	3
---	---	---	---	---	---	---

row indices

0	2	2	5	7
---	---	---	---	---


# Compressed Sparse Row (CSR) Format

---


- **M** number of rows in the matrix
- **N** number of columns in the matrix
- **S** sparsity level  $[0 - 1]$ , 1 being fully-dense
- **Space Requirements**
  - Dense Representation  $M*N$
  - Sparse representation with CSR  $2MNS + M + I$
- CSR only saves space if  $S < (I - N - 1) / 2$ 
  - Otherwise, indexes saving introduce overhead

# Parallel SpMV CSR

- Assign a thread to each row of the matrix
  - The thread loops through the nonzero elements of its row to perform the dot product
  - A single thread traverses each row
  - So, each thread will write to a distinct output value



Logical view



Physical view



# Parallel SpMV CSR - Considerations

---

- CSR is **space efficient**
  - Higher than COO as the row pointers vector is smaller than the row indexes one
- CSR is **not flexible**
  - Adding non-zero elements to the matrix is expensive
    - As it involves shifting elements
- CSR is **easy to access by row**
  - Easy to find non-zero elements in a row
    - The same does not hold for columns
  - This is true with big matrices with a lot of rows

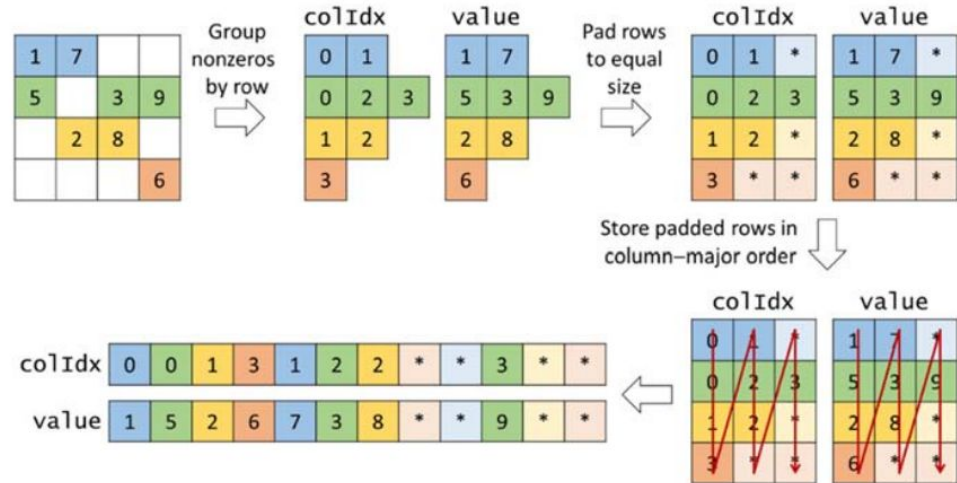
# Parallel SpMV CSR - Considerations

---

- CSR is not using the memory bandwidth efficiently
  - Consecutive threads will access memory locations far apart from each other
  - No coalesced memory access
- CSR potential flow divergence in warps
  - The number of iterations that are taken by a thread in the dot product loop depends on the number of nonzero elements in the row that is assigned to the thread
- NO Atomic operation is required

# ELL Format

- The name came from the sparse matrix package in ELLPACK
  - A package for solving elliptic boundary value problems
  - Relies on the maximum number of non-zeros elements per row
  - Usage of paddings elements



# ELL Format



---

- **M** number of rows in the matrix
- **N** number of columns in the matrix
- **K** number of nonzero entries in the densest row
- **S** sparsity level  $[0 - 1]$ , 1 being fully-dense
- **Space Requirements**
  - Dense Representation  $M*N$
  - Sparse representation with ELL  $2MK$
- ELL only saves space if  $K < N / 2$



# Parallel SpMV ELL

- Each thread is assigned to a different row of the matrix
  - A dot product loop goes through the non-zero elements of each row
  - Each thread iterates only through the non-zeros in its assigned row
    - Thanks to the maximum number of non-zero elements per row



# Parallel SpMV ELL - Considerations

---

- ELL is less space efficient than CSR
  - Due to the presence of padding elements
  - The compression rate depends of the maximum number on non-zeros elements per row
- ELL is more flexible than CSR
  - A non-zero element can be added by replacing a padding element
- ELL is easy to access
  - Given a value  $i$ , we can get its column value and row easily
  - We can get parallelize by rows and by non-zero elements

# Parallel SpMV ELL - Considerations


---

- ELL uses the memory bandwidth efficiently
  - Consecutive threads perform coalesced access to memory
- ELL computation is imbalanced
  - As each thread iterates over the non-zero elements of its row
- NO Atomic operation is required



# The Coordinate (COO) Format

- Store both the column index and row index for every non-zero
  - No ordering is required for indexes
  - Overhead because all indexes are stored
- COO and CSR format differs since CSR replaces the row indexes array with row pointers that store the starting offset of each row's nonzeros in the other arrays







# The Coordinate (COO) Format

---

- **M** number of rows in the matrix
- **N** number of columns in the matrix
- **S** sparsity level  $[0 - 1]$ , 1 being fully-dense
- **Space Requirements**
  - Dense Representation  $MN$
  - Sparse representation with COO  $3MNS$
- COO only saves space if  $S < 1 / 3$

# Parallel SpMV COO

- Each thread is responsible for the computation of a non-zero element in the matrix



# Parallel SpMV COO - Considerations

---

- COO is less space efficient than CSR
  - It requires more space due to the repeated indexes
- COO is flexible
  - Non-zero elements can be pushed back into the arrays
- COO accessibility varies
  - We can process elements in any order
  - We cannot access all non-zeros elements by row and by column easily

# Parallel SpMV COO - Considerations

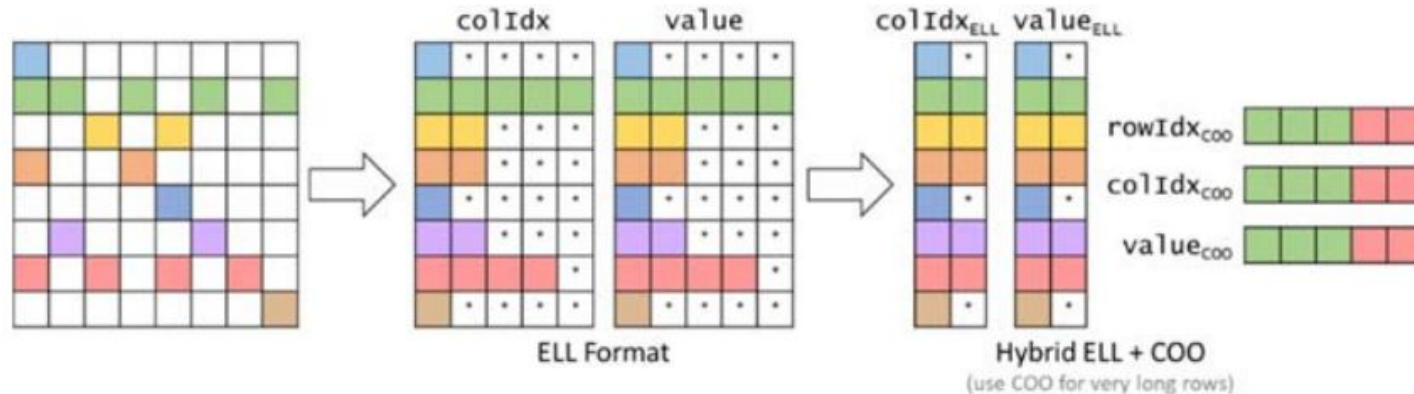
---

- COO uses the memory bandwidth efficiently
  - Threads access the indexes array in a coalesced way
- COO computation is balanced
  - All threads are responsible for the same amount of work
- Atomic operations ARE required



# Hybrid ELL/COO Format

- ELL format suffers from rows with a large amount of non-zero elements
  - Mostly in terms of space efficiency and control divergence
- Place non-zeros from the densest rows in a COO sparse matrix, leading to a more efficient ELL representation for the remainder
  - Each element will be stored in the ELL or the COO matrix, not both



# Parallel SpMV Hybrid ELL/COO Format

---

- Perform the SpMV / ELL in parallel on the GPU
  - Rows have a similar density, thus the execution is more efficient
- Perform the SpMV / COO sequentially on the CPU
  - CPUs better handle the COO access irregularity
- The increased overhead has to be justified
  - If the SpMV is computed only once probably, the execution time will increase
  - If the SpMV is done inside an iterative solver, the execution time is likely to decrease

# Parallel SpMV Hybrid ELL/COO - Considerations

---

- Hybrid ELL/COO is more space efficient than ELL
  - It reduces padding
- Hybrid ELL/COO is more flexible than ELL
  - The non-zero elements can be added
    - For the ELL part, if the element in the column is available
    - Otherwise, it is pushed back in the COO part
- Hybrid ELL/COO accessibility is reduced wrt to ELL
  - If the elements are not available in the ELL part, a search on the COO part is required

# Parallel SpMV Hybrid ELL/COO - Considerations

---

- Hybrid ELL/COO uses the memory bandwidth efficiently
  - Both the ELL and COO parts are accessed in a coalesced way
- Hybrid ELL/COO computation is more balanced than ELL
  - Less padding, less divergence





# Storage Requirements Summary


---

- **M** number of rows in the matrix
- **N** number of columns in the matrix
- **K** number of nonzero entries in the densest row
- **S** sparsity level [0 -1], 1 being fully-dense

Format	Storage Requirements(words)
Dense	$MN$
Compressed Sparse Row (CSR)	$2MNS + M + 1$
ELL	$2MK$
Coordinate (COO)	$3MNS$
Hybrid ELL / COO (HYB)	$2MK < 3MNS$


# Jagged Diagonal Storage (JDS) Format

- Group similarly dense rows into evenly sized partitions and represent each section independently using either CSR or ELL
  - Sort rows by density
- Need to store the original rows of the sorted rows
  - To preserve re-constructability



# Parallel SpMV Jagged Diagonal Storage (JDS) Format

- Each thread is assigned to a row of the matrix and iterates through the nonzeros of that row
  - performing the dot product along the way
- You can also partition the matrix into sections of similar rows
  - These partitions can be compressed using ELL



# Parallel SpMV JDS - Considerations

---

- JDS is more space efficient than ELL
  - It avoids and reduces padding
- JDS is not flexible
  - Add new elements is complex
    - It may require new sorting
- JDS accessibility is similar to CSR
  - It allows us to access, given a row index, the nonzero elements of that row
  - It does not make it easy to access, given a nonzero, the row index and column index of that nonzero

# Parallel SpMV Hybrid ELL/COO - Considerations

---

- JDS uses the memory bandwidth efficiently
  - JDS allows memory access to be coalesced
- JDS computation is balanced
  - Rows are sorted so that threads in the same warp execute similar work

# Other Formats

---

- **Diagonal (DIA)**
  - Stores only a sparse set of dense diagonal vectors
  - For each diagonal, the offset from the main diagonal is stored
- **Packet (PKT)**
  - Reorders rows and columns to concentrate nonzeros into roughly diagonal submatrices
  - This improves cache performance as nearby rows access nearby x elements
- **Dictionary of Keys (DOK)**
  - Matrix is stored as a map from (row,column) index pairs to values
  - This can be useful for building or querying a sparse matrix, but iteration is slow
- **Compressed Sparse Column (CSC)**
  - Like CSR, but stores a dense set of sparse column vectors
  - Useful for when column sparsity is much more regular than row sparsity
- **Blocked CSR**
  - The matrix is divided into blocks stored using CSR with the indices of the upper left corner
  - Useful for block-sparse matrices
- **Additional Hybrid Methods**
  - For example, DIA is very inefficient when there are a small number of mostly-dense diagonals, but a few additional sparse entries
  - In this case, a hybrid DIA / COO or DIA / CSR representation can be used

# Conclusions

---

- In general, the FLOPS ratings that are achieved by either CPUs or GPUs are much lower for sparse matrix computation than for dense matrix computation
  - In SpMV computation, there is no data reuse in the sparse matrix
- The  $OP/B$  is essentially  $0.25$ , limiting the achievable FLOPS rate to a small fraction of the peak performance
- The various formats are important for both CPUs and GPUs since both are limited by memory bandwidth

# Other Libraries

---

- **cuSPARSE**

- cuSPARSE is an Nvidia library implemented in set of basic linear algebra subroutines (BLAS)
- Reference available [here](#)

- **CUSP**

- CUSP is an open-source project
- It is focused on solving linear systems, providing a number of conjugate-gradient solvers
- Repository available [here](#)



# Credits

---

- [CSE 599 I](#) Accelerated Computing - Programming GPUS - Parallel Pattern: Sparse Matrices, Author *Tanner Schmidt*



# Thank you for your attention!

**Gianmarco Accordi**  
*[gianmarco.accordi@polimi.it](mailto:gianmarco.accordi@polimi.it)*